

1 Memory Hierarchy

Processor (registers)	Fast, small, expensive
Cache	↓
Main Memory	↓
Disk	Slow, large, cheap

- **Hit**: data appears in some block in the upper levels of memory (closer to the processor)
 - **Hit Rate**: the fraction of memory access found in the upper level
 - **Hit Time**: the time to access the upper level
= memory access time + time to determine hit/miss
- **Miss**: data needs to be retrieved from a block in the lower level of memory.
 - **Miss Rate**: $1 - (\text{hit rate})$
 - **Miss Penalty** = time to replace block in upper level + time to deliver block to processor
- **average access time**
= hit time \times hit rate + miss penalty \times miss rate
- **average memory access time (AMAT)**
= hit time + miss penalty \times miss rate

2 Cache

2.1 Direct-Mapped Cache

- Each memory address is associated with one possible *block* within the cache
- **Block**: the unit of transfer between cache and memory
- Address in Direct-Mapped Cache

t	a	g	idx	o
---	---	---	-----	---

 - **tag**: to check if have correct block
 - **index**: to select block (which "row")
 - **byte offset** within block
- larger cache blocks take better advantage of spacial locality
- every block has a *valid bit* determines whether anything stored in that row

2.1.1 Example

16KB of data in direct-mapped cache with 4 word blocks. Determine size of tag, index and offset fields (given 32-bit architecture)

- **index**: $\# \text{ blocks} = 16 \text{ KB} / 2^4 = 2^{10} \text{ B} \rightarrow 10 \text{ bits}$
- **byte offset**: $4 \text{ words} = 16 \text{ bytes} = 2^4 \text{ bytes} \rightarrow 4 \text{ bits}$
- **tag**: $32 - 4 - 10 = 18 \rightarrow 18 \text{ bits}$

2.1.2 Impact of Large Block Size

- **benefit**: reduce miss rate (takes advantage of spacial locality)
- **drawbacks**: larger miss penalty (takes longer to load new block from next level) and possibly a higher miss rate

2.2 Fully Associative Cache

- no rows, basically a hash map
- **benefits**: no conflicts among different memory addresses
- **drawbacks**: need hardware comparator for every single entry: this is infeasible

2.3 N-Way Set Associative Cache

- memory address fields: tag (same as before), offset (same as before), index (points us to correct row (called a set in this case))

- cache is direct-mapped with respect to sets, each set is set associative
- has best of direct-mapped and fully associative

2.4 Examples: Number of Bits Needed

64 KB of data and one word blocks

$$64 \text{ KB} = 16 \text{ K words} = 2^{14} \text{ words} = 2^{14} \text{ blocks}$$

• direct-mapped cache

block size = 4 bytes \rightarrow offset size = 2 bits

$\# \text{ sets} = \# \text{ blocks} = 2^{14} \rightarrow$ index size = 14 bits

tag size = $32 - 14 - 2 = 16 \text{ bits}$

bits/block = data + tag + valid = $32 + 16 + 1 = 49$

bits/cache = $\# \text{ blocks} \times \text{bit/block} = 2^{14} \times 49 = 98 \text{ KB}$

• 4-way set associative

block size = 4 bytes \rightarrow offset size = 2 bits

$\# \text{ sets} = \# \text{ blocks}/4 = 2^{14}/4 = 2^{12} \rightarrow$ index size = 12 bits

tag size = $32 - 12 - 2 = 18 \text{ bits}$

bits/block = data + tag + valid = $32 + 18 + 1 = 51$

bits/cache = $\# \text{ blocks} \times \text{bit/block} = 2^{14} \times 51 = 102 \text{ KB}$

• increase associativity \rightarrow increase bits in cache

• 8 word blocks

$64 \text{ KB} = 2^{14} \text{ words} = 2^{14}/8 \text{ blocks} = 2^{11} \text{ blocks}$

block size = 4 words = 32 bytes \rightarrow offset size = 5 bits

$\# \text{ sets} = \# \text{ blocks} = 2^{11} \rightarrow$ index size = 11 bits

tag size = address - index - offset = $32 - 11 - 5 = 16 \text{ bits}$

bits/block = data + tag + valid = $32 + 16 + 1 = 49 \text{ bits}$

bits/cache = $\# \text{ blocks} \times \text{bit/block} = 2^{11} \times 49 = 68.25 \text{ KB}$

• increase block size \rightarrow decrease bits in cache

2.5 Accessing a Cache

- **cache hit**: cache block is valid and contains proper address, read the desired word
- **cache miss**: nothing in cache in appropriate block, fetch from memory
- **cache miss, block replacement**: wrong data is in cache at appropriate block, so discard it and fetch desired data from memory
- **Types of Cache Misses**
 - Compulsory Miss: occur when program first starts
 - Conflict Misses: occurs when two addresses map to the same cache location (problem in direct-mapped cache)
 - Capacity Misses: miss that occurs because the cache has limited size (problem in fully associative caches)

2.6 Block Replacement

- if there are any locations with valid bit of 0, usually write over that
- if all possible locations already have valid block, choose replacement strategy
 - random (simple to implement)
 - least-recently used (LRU) (expensive)

2.7 Multit-Level Cache Hierarchy

Given 2 levels of cache:

$$\text{AMAT} = \text{L1 hit time} + \text{L1 miss rate} \times \text{L1 miss penalty}$$

$$\text{L1 miss penalty} = \text{L2 hit time} + \text{L2 miss rate} \times \text{L2 miss penalty}$$

3 Examples

3.1 Average Memory Address Time (AMAT)

Assume: hit time = 1 cycle, miss rate = 5%, miss penalty = 20 cycles

$$\text{AMAT} = 1 + 0.05 \times 20 = 1 + 1\text{cycles} = 2\text{cycles}$$