

1 Overview

- **tuple**: an ordered and named collection of values
Ex: ('Mehta, Dinesh', 'CSCI 406', 'ALGORITHMS')
- **attributes**: formally, the attributes are a set
Ex: ('instructor', 'course-id', 'title')
- **domain**: each attribute has associated domain from which values come
- **relation schema**:
 - A relation schema R is denoted as $R(A_1, A_2, \dots, A_n)$.
 - Has **degree** n
 - Attributes A_i
 - Each attribute has associated domain: $Dom(A_i)$
- **relation or relation state**:
 - A relation is a set of tuples $r(R)$ conforming to the relation schema: for any tuple $t \in r$

$$t = (v_1, v_2, \dots, v_n) : v_i \in Dom(A_i)$$

- In other words: $r(R) \subseteq (dom(A_1) \times dom(A_2) \times \dots \times dom(A_n))$
- **In actuality** real RDBMS
 - * tables may contain duplicates
 - * no ordering of tuples
- **NULL**
A "value" that can exist in a relational database representing
 - a value is unknown
 - value is irrelevant
 - value is missing

problems with NULL

- NULL values cannot be compared

$$a = b, a \text{ is NULL}, b \text{ is NULL} \rightarrow \text{false}$$

instead we have ISNULL operator

- **constraints**
 - on the relation schema
 - between relation schemas
- **superkey** - A superkey of R is a subset of R 's attributes such that no relation $r(R)$ may contain tuples with exactly the same values for attributes in the superkey.
 - any superset of a superkey is also a superkey
 - every relation (in theory) has at least one superkey - the set of all attributes of the schema
- **Key** - A superkey s.t. no attribute can be removed from the key without destroying the superkey property: a *minimal* superkey
- **Candidate Keys** - All the keys of a schema are called *candidate keys*
- **Primary Key** - By convention, we choose one candidate key as the *primary key*
 - we tend to prefer "smaller" keys (e.g. *crn*)
 - in practice in SQL DB:
 - * primary keys impose uniqueness constraint on data
 - * other candidate keys \rightarrow impose uniqueness constraint
- **Referential Integrity/Foreign Key Constraint** Constraint on the relationship between relations. A *foreign key* is a subset of attributes of a relation schema with the property that its values are either NULL or existing in the specified attributes of a referenced table.

2 Relational Theory

- each relation is a **set** of tuples
- tuple = set of values associated with named attributes
- **Relational Algebra**
 - useful operations that can be applied to relations
 - resulting algebra

3 Unary Operations $-2 + 1$

3.1 Selection

- Chooses subsets of tuples from a relation according to some condition (**WHERE**)
- Letting R represent some relation: $\sigma_{\text{condition}}(R)$
- E.g. $\sigma_{\text{course_id} = 'CSCI 403'}(\text{mines_courses})$
- Can use AND, OR, NOT, etc.
- Properties of σ
 - degree of relation ($\#$ of attributes) is the same as the original
 - the number of tuples in the result $\leq \#$ tuples in original
 - commutative $\sigma_a(\sigma_b(R)) = \sigma_b(\sigma_a(R))$
 - any sequence of σ operations can be replaced with a single σ (using AND)

3.2 Projection

- chooses attributes from set of attributes in a relation
- parallel operation in SQL: **SELECT**
- Notated:

$$\pi_{\text{attr}_1, \text{attr}_2, \dots}(R)$$

- Example:

$$\pi_{\text{instructor}, \text{course_id}}(\text{mines_courses})$$

- Properties of π
 - $\#$ of tuples is always $\leq \#$ tuples in original
 - **not** commutative: rather

$$\pi_{\text{list}_1}(\pi_{\text{list}_2}(R)) = \pi_{\text{list}_1}(R)$$

if $\text{list}_1 \subset \text{list}_2$
else, ill formed

3.3 Renaming - ρ

- parallel operation in SQL: **AS**
- Notation:

$$\rho_{S(B_1, B_2, \dots)}(R)$$

- Example: Let table X have attributes (a, b, c) .

$$\rho_{Y(d, e, f)}(X)$$

```
SELECT a AS d,
       b AS e,
       c AS f
FROM X AS Y;
```

3.4 Sequences of Operations

3.4.1 Representations

1. nesting: example

$$\rho_{(\text{name}, \text{courseid})}(\pi_{\text{instructor}, \text{courseid}}(\sigma_{\text{department} = 'CS'}(\text{mines_cour.})))$$

```
SELECT instructor AS name
       course_id
FROM   mines_courses
WHERE  department = 'CS';
```

2. sequence of named relations

$$R_1 = \sigma_{department='CS'}(mines_courses)$$

$$R_2 = \pi_{instructor, course_id}(R_1)$$

$$R_3 = \rho_{(name, course_id)}(R_2)$$

4 Binary Operations

4.1 Set Operators

- $A \cup B$ - union
- $A \cap B$ - intersection
- $A - B$ - difference

4.1.1 Properties of Set Operators

- Union, intersection are commutative and associative
- Set difference is not commutative or associative

4.2 Cartesian Product and Joins

- $A \times B$ - pairs every tuple from A with every tuple from B
- If A has m attributes, B has n , $A \times B$ has $m + n$ attributes
 $|A \times B| = |A| \times |B|$.
- Typical Usage:

$$\sigma_{condition}(A \times B) \equiv A \bowtie_{condition} B$$

- Example: $mines_courses \bowtie_{instructor=name} mines_eecs_faculty$

4.3 Theta Joins

- When:
 - we have $A \bowtie_{cond} B$
 - condition is of general form $cond_1 AND cond_2 AND \dots$
 - each $cond_n$ $A_i \Theta B_j$, $A_i \in A$, $B_j \in B$

Then we call $A \bowtie_{cond}$ a *theta join*: $A \bowtie_{\Theta} B$

- Further, when Θ is $=$, then $A \bowtie_{\Theta} B$ is called an *equijoin* which implies that there is some duplicate data column.
- If $A \bowtie_{cond} B$ is an equijoin and condition equates attributes of A with attributes of B of the **same name**, then the join is a *natural join*.

$A * B$ - books notation

$A \bowtie B$ - other people's notation

Example: both `mines_courses` and `mines_courses_meetings` have `crn` field. Then $mines_courses * mines_courses_meetings$.

(Note, this automatically projects away duplicate column(s).)

5 Completeness

You can demonstrate that: $\sigma, \pi, \cup, -, \times$ is a complete set of operators for relational algebra.

6 Odds and Ends

- Division: \div ~inverse of \times (no mirror in SQL)
- Aggregates and grouping: not part of basic algebra